

Robotic Simulation Incident Analysis

A failed robot run gets a root cause, a fix, and a ticket. Then it re-runs to prove the fix held. One command, 13 seconds.

0% → 100%

ascent success, before → after fix

5 s · \$0.0003

per diagnosis, Gemma 4 31B via Respan

1 Run fails

G1 climbs the fixed line with balance assist at 0.5. Pelvis height decays from 0.63 m, fall at step 126. Every step logged.

2 Diagnose

Telemetry and tunables go to Gemma 4 31B through the Respan gateway. Back comes root cause, evidence, and a fix as JSON.

3 Fix and re-run

Assist 0.5 → 1.0 on the live environment, same seed. Robot reaches 1.53 m. Success, not a guess.

4 File the issue

Reviewer on a Lambda GPU confirms the diagnosis, then Nango files the GitHub issue with evidence and the verdict.

HOW EACH SPONSOR IS USED

■ Respan

Gateway for every model call. Traces are the training set, evals gate the student model.

■ Gemma

Gemma 4 31B is the analyst today. Gemma 4 E4B is the student to distill into.

■ Lambda

Serves the reviewer model (Gemma 3 12B on a GH200) that verifies every diagnosis. Fine-tune next.

■ Nango

GitHub and Slack tools with auth handled. The agent files the ticket.